

KingCountyHousing

January 1, 2018

1 Create the dataframe

```
In [8]: import pandas as pd
        housing_data = pd.read_csv("Data/KingCountyHousing.csv")
        # housing_data.head()
```

Here we have 21 Columns with 19 Features, and 21613 Observations.

2 Fix the date

```
In [13]: import datetime
        current_year = datetime.datetime.now().year
        housing_data["age_of_house"] = current_year - pd.to_datetime(
            housing_data["date"]).dt.year
```

We've now added a column for the age of the home, which will help us analyze the pricing.

```
In [14]: housing_data.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 21613 entries, 0 to 21612
Data columns (total 22 columns):
id                21613 non-null int64
date              21613 non-null object
price             21613 non-null float64
bedrooms          21613 non-null int64
bathrooms         21613 non-null float64
sqft_living       21613 non-null int64
sqft_lot          21613 non-null int64
floors            21613 non-null float64
waterfront        21613 non-null int64
view              21613 non-null int64
condition         21613 non-null int64
grade             21613 non-null int64
sqft_above        21613 non-null int64
sqft_basement     21613 non-null int64
yr_built          21613 non-null int64
```

```
yr_renovated      21613 non-null int64
zipcode           21613 non-null int64
lat               21613 non-null float64
long              21613 non-null float64
sqft_living15     21613 non-null int64
sqft_lot15        21613 non-null int64
age_of_house      21613 non-null int64
dtypes: float64(5), int64(16), object(1)
memory usage: 3.6+ MB
```

```
In [15]: housing_data.columns
```

```
Out[15]: Index(['id', 'date', 'price', 'bedrooms', 'bathrooms', 'sqft_living',
               'sqft_lot', 'floors', 'waterfront', 'view', 'condition', 'grade',
               'sqft_above', 'sqft_basement', 'yr_built', 'yr_renovated', 'zipcode',
               'lat', 'long', 'sqft_living15', 'sqft_lot15', 'age_of_house'],
              dtype='object')
```

3 Select features and dependent variable


```
In [65]: feature_cols = [
           u'age_of_house',
           u'bedrooms',
           u'bathrooms',
           u'sqft_living',
           u'sqft_lot',
           u'floors',
           u'waterfront',
           u'view',
           u'condition',
           u'grade',
           u'sqft_above',
           u'sqft_basement',
           u'yr_built',
           u'yr_renovated',
           u'zipcode',
           u'lat',
           u'long',]
x = housing_data[feature_cols]
y = housing_data["price"]
```

4 Visualize the features against the dependent variable

```
In [19]: import matplotlib.pyplot as plt
         %matplotlib inline
```


```
plt.title("Age of House")
plt.scatter(housing_data["age_of_house"],housing_data["price"])
```

Out[19]: <matplotlib.collections.PathCollection at 0x285dca962b0>




```
In [20]: plt.title("bedrooms")
plt.scatter(housing_data["bedrooms"],housing_data["price"])
```

Out[20]: <matplotlib.collections.PathCollection at 0x285dcadf320>




```
In [22]: plt.title("bathrooms")
         plt.scatter(housing_data["bathrooms"],housing_data["price"])
```

```
Out[22]: <matplotlib.collections.PathCollection at 0x285dcb53400>
```




```
In [23]: plt.title("sqft_living")
         plt.scatter(housing_data["sqft_living"],housing_data["price"])
```

```
Out[23]: <matplotlib.collections.PathCollection at 0x285dcbb9908>
```




```
In [24]: plt.title("sqft_lot")
         plt.scatter(housing_data["sqft_lot"],housing_data["price"])
```

```
Out[24]: <matplotlib.collections.PathCollection at 0x285dcc20dd8>
```




```
In [25]: plt.title("floors")
         plt.scatter(housing_data["floors"],housing_data["price"])
```

```
Out[25]: <matplotlib.collections.PathCollection at 0x285dcf58ba8>
```




```
In [26]: plt.title("waterfront")
         plt.scatter(housing_data["waterfront"],housing_data["price"])
```

```
Out[26]: <matplotlib.collections.PathCollection at 0x285dcf96ba8>
```




```
In [27]: plt.title("view")
         plt.scatter(housing_data["view"],housing_data["price"])
```

```
Out[27]: <matplotlib.collections.PathCollection at 0x285dd028080>
```




```
In [28]: plt.title("condition")
         plt.scatter(housing_data["condition"],housing_data["price"])
```

```
Out[28]: <matplotlib.collections.PathCollection at 0x285dd08d8d0>
```





```
In [29]: plt.title("grade")
plt.scatter(housing_data["grade"],housing_data["price"])
```

```
Out[29]: <matplotlib.collections.PathCollection at 0x285dd0fe198>
```



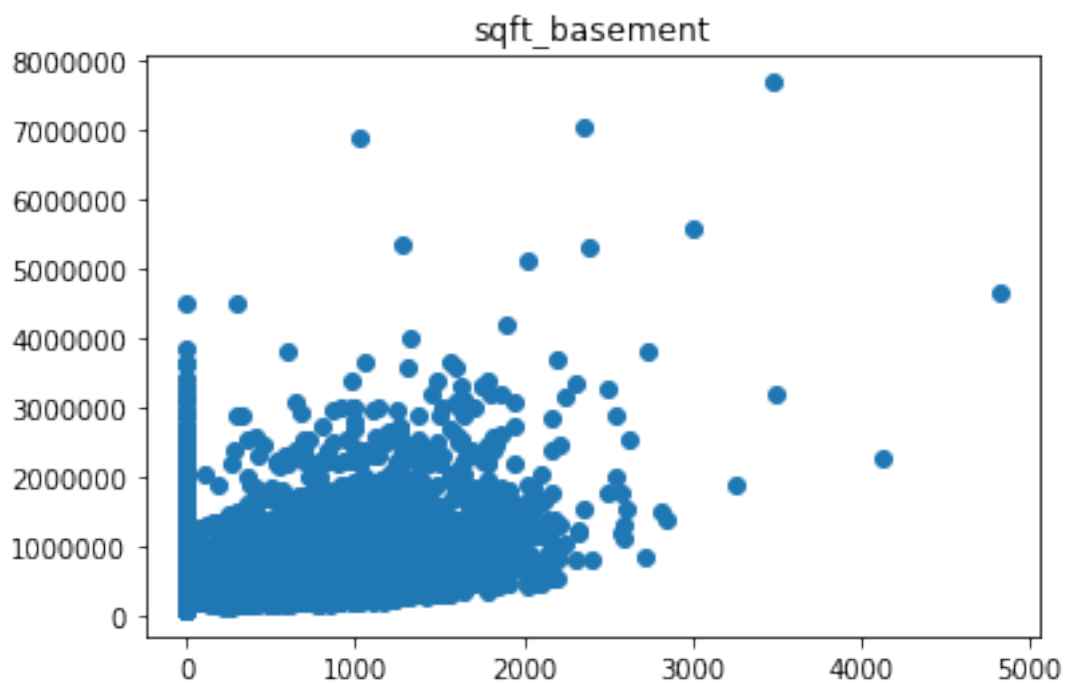
```
In [30]: plt.title("sqft_above")
plt.scatter(housing_data["sqft_above"],housing_data["price"])
```

```
Out[30]: <matplotlib.collections.PathCollection at 0x285dd15e860>
```




```
In [31]: plt.title("sqft_basement")
         plt.scatter(housing_data["sqft_basement"],housing_data["price"])
```

```
Out[31]: <matplotlib.collections.PathCollection at 0x285dd1bce48>
```




```
In [32]: plt.title("yr_built")
plt.scatter(housing_data["yr_built"],housing_data["price"])
```

```
Out[32]: <matplotlib.collections.PathCollection at 0x285dd2267b8>
```




```
In [33]: plt.title("yr_renovated")
plt.scatter(housing_data["yr_renovated"],housing_data["price"])
```

```
Out[33]: <matplotlib.collections.PathCollection at 0x285dd28d6a0>
```




```
In [66]: plt.title("zipcode")
         plt.scatter(housing_data["zipcode"],housing_data["price"])
```

```
Out[66]: <matplotlib.collections.PathCollection at 0x285e5dac390>
```




```
In [67]: plt.title("lat")
         plt.scatter(housing_data["lat"],housing_data["price"])

Out[67]: <matplotlib.collections.PathCollection at 0x285e5e00160>
```



```
In [68]: plt.title("long")
         plt.scatter(housing_data["long"],housing_data["price"])

Out[68]: <matplotlib.collections.PathCollection at 0x285e5e61588>
```



5 Split the training and test data

```
In [69]: from sklearn.model_selection import train_test_split
         x_train,x_test,y_train,y_test = train_test_split(x, y, test_size=.2, random_state=3)
```

6 Fitting the model to the training set

```
In [70]: from sklearn.linear_model import LinearRegression
         regressor = LinearRegression()
         regressor.fit(x_train, y_train)
```

```
         accuracy = regressor.score(x_test, y_test)
         print(accuracy)
```

0.709354281043

```
In [78]: # Building the optimal model using Backward Elimination
         def backwardElim(X_opt, SL):
             regressor_OLS = sm.OLS(endog = y, exog = X_opt).fit()
             for i in range(np.size(X_opt,1)):
                 if regressor_OLS.pvalues[i] > SL:
                     if regressor_OLS.pvalues[i] == max(regressor_OLS.pvalues):
                         print(regressor_OLS.summary())
```

```

        print("removing: "+str(i)+" , with P val: "+str(regressor_OLS.pvalues[i])
        return backwardElim(np.delete(X_opt, i, axis=1), SL)
    return X_opt

```

```

In [92]: import statsmodels.formula.api as sm
        X = np.append(arr =np.ones((21613,1)).astype(int), values = x, axis = 1)
        X_opt = X[:, [0,1,2,3,4,5,6,7,8,9,10,11,12,13,14,15,16,17]]
        SL = 0.05

        X_opt = backwardElim(X_opt, SL)
        regressor_OLS = sm.OLS(endog = y, exog = X_opt).fit()

```

OLS Regression Results

```

=====
Dep. Variable:          price    R-squared:                0.700
Model:                  OLS      Adj. R-squared:           0.700
Method:                 Least Squares    F-statistic:        3155.
Date:                  Mon, 01 Jan 2018    Prob (F-statistic):    0.00
Time:                  15:03:35    Log-Likelihood:       -2.9458e+05
No. Observations:      21613    AIC:                  5.892e+05
Df Residuals:          21596    BIC:                  5.893e+05
Df Model:               16
Covariance Type:       nonrobust
=====

```

| | coef | std err | t | P> t | [0.025 | 0.975] |
|-------|------------|----------|---------|-------|-----------|-----------|
| const | 9.362e+06 | 2.88e+06 | 3.254 | 0.001 | 3.72e+06 | 1.5e+07 |
| x1 | -3.04e+04 | 2932.859 | -10.365 | 0.000 | -3.61e+04 | -2.47e+04 |
| x2 | -3.569e+04 | 1887.879 | -18.905 | 0.000 | -3.94e+04 | -3.2e+04 |
| x3 | 4.11e+04 | 3247.663 | 12.656 | 0.000 | 3.47e+04 | 4.75e+04 |
| x4 | 114.7710 | 2.127 | 53.953 | 0.000 | 110.601 | 118.941 |
| x5 | -0.0551 | 0.035 | -1.589 | 0.112 | -0.123 | 0.013 |
| x6 | 5506.8689 | 3567.325 | 1.544 | 0.123 | -1485.352 | 1.25e+04 |
| x7 | 5.797e+05 | 1.73e+04 | 33.442 | 0.000 | 5.46e+05 | 6.14e+05 |
| x8 | 5.452e+04 | 2111.286 | 25.821 | 0.000 | 5.04e+04 | 5.87e+04 |
| x9 | 2.692e+04 | 2350.233 | 11.452 | 0.000 | 2.23e+04 | 3.15e+04 |
| x10 | 1.001e+05 | 2060.886 | 48.580 | 0.000 | 9.61e+04 | 1.04e+05 |
| x11 | 74.5709 | 2.136 | 34.905 | 0.000 | 70.383 | 78.758 |
| x12 | 40.2011 | 2.643 | 15.212 | 0.000 | 35.021 | 45.381 |
| x13 | -2637.2436 | 72.548 | -36.352 | 0.000 | -2779.442 | -2495.045 |
| x14 | 19.3364 | 3.648 | 5.301 | 0.000 | 12.187 | 26.486 |
| x15 | -603.0424 | 32.804 | -18.383 | 0.000 | -667.341 | -538.743 |
| x16 | 6.098e+05 | 1.07e+04 | 56.935 | 0.000 | 5.89e+05 | 6.31e+05 |
| x17 | -2.078e+05 | 1.29e+04 | -16.090 | 0.000 | -2.33e+05 | -1.82e+05 |

```

=====
Omnibus:                18172.766    Durbin-Watson:           1.992
Prob(Omnibus):           0.000    Jarque-Bera (JB):        1785086.012
Skew:                    3.509    Prob(JB):                0.00
=====

```

Kurtosis: 46.966 Cond. No. 8.00e+16

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
 [2] The smallest eigenvalue is 3.34e-20. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.
 removing: 6, with P val: 0.122676448789

OLS Regression Results

```

=====
Dep. Variable:          price    R-squared:                0.700
Model:                  OLS      Adj. R-squared:            0.700
Method:                 Least Squares    F-statistic:          3365.
Date:                  Mon, 01 Jan 2018    Prob (F-statistic):      0.00
Time:                  15:03:35    Log-Likelihood:         -2.9458e+05
No. Observations:      21613    AIC:                    5.892e+05
Df Residuals:          21597    BIC:                    5.893e+05
Df Model:               15
Covariance Type:       nonrobust
=====
  
```

| | coef | std err | t | P> t | [0.025 | 0.975] |
|-------|------------|----------|---------|-------|-----------|-----------|
| const | 8.47e+06 | 2.82e+06 | 3.005 | 0.003 | 2.95e+06 | 1.4e+07 |
| x1 | -3.032e+04 | 2932.545 | -10.340 | 0.000 | -3.61e+04 | -2.46e+04 |
| x2 | -3.576e+04 | 1887.390 | -18.947 | 0.000 | -3.95e+04 | -3.21e+04 |
| x3 | 4.244e+04 | 3130.401 | 13.557 | 0.000 | 3.63e+04 | 4.86e+04 |
| x4 | 114.2382 | 2.099 | 54.424 | 0.000 | 110.124 | 118.353 |
| x5 | -0.0583 | 0.035 | -1.684 | 0.092 | -0.126 | 0.010 |
| x6 | 5.799e+05 | 1.73e+04 | 33.451 | 0.000 | 5.46e+05 | 6.14e+05 |
| x7 | 5.454e+04 | 2111.312 | 25.830 | 0.000 | 5.04e+04 | 5.87e+04 |
| x8 | 2.672e+04 | 2346.876 | 11.385 | 0.000 | 2.21e+04 | 3.13e+04 |
| x9 | 1.003e+05 | 2057.307 | 48.757 | 0.000 | 9.63e+04 | 1.04e+05 |
| x10 | 75.7444 | 1.996 | 37.939 | 0.000 | 71.831 | 79.658 |
| x11 | 38.4927 | 2.400 | 16.038 | 0.000 | 33.788 | 43.197 |
| x12 | -2612.3789 | 70.739 | -36.930 | 0.000 | -2751.033 | -2473.725 |
| x13 | 19.6660 | 3.641 | 5.401 | 0.000 | 12.529 | 26.803 |
| x14 | -597.7268 | 32.624 | -18.322 | 0.000 | -661.672 | -533.781 |
| x15 | 6.111e+05 | 1.07e+04 | 57.230 | 0.000 | 5.9e+05 | 6.32e+05 |
| x16 | -2.1e+05 | 1.28e+04 | -16.353 | 0.000 | -2.35e+05 | -1.85e+05 |

```

=====
Omnibus:                18138.075    Durbin-Watson:          1.991
Prob(Omnibus):           0.000    Jarque-Bera (JB):       1772812.098
Skew:                    3.499    Prob(JB):                0.00
Kurtosis:                46.814    Cond. No.                7.42e+16
=====
  
```

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The smallest eigenvalue is 3.89e-20. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.
removing: 5, with P val: 0.0922781409562

In [94]: regressor_OLS.summary()

Out[94]: <class 'statsmodels.iolib.summary.Summary'>
"""

```

                                OLS Regression Results
=====
Dep. Variable:                  price    R-squared:                  0.700
Model:                            OLS    Adj. R-squared:              0.700
Method:                 Least Squares    F-statistic:                 3604.
Date:                   Mon, 01 Jan 2018    Prob (F-statistic):          0.00
Time:                   15:04:38    Log-Likelihood:             -2.9458e+05
No. Observations:          21613    AIC:                       5.892e+05
Df Residuals:              21598    BIC:                       5.893e+05
Df Model:                   14
Covariance Type:            nonrobust
=====

```

| | coef | std err | t | P> t | [0.025 | 0.975] |
|-------|------------|----------|---------|-------|-----------|-----------|
| const | 7.969e+06 | 2.8e+06 | 2.843 | 0.004 | 2.48e+06 | 1.35e+07 |
| x1 | -3.029e+04 | 2932.588 | -10.328 | 0.000 | -3.6e+04 | -2.45e+04 |
| x2 | -3.551e+04 | 1881.707 | -18.873 | 0.000 | -3.92e+04 | -3.18e+04 |
| x3 | 4.259e+04 | 3129.276 | 13.609 | 0.000 | 3.65e+04 | 4.87e+04 |
| x4 | 113.8393 | 2.086 | 54.579 | 0.000 | 109.751 | 117.928 |
| x5 | 5.802e+05 | 1.73e+04 | 33.472 | 0.000 | 5.46e+05 | 6.14e+05 |
| x6 | 5.438e+04 | 2109.328 | 25.780 | 0.000 | 5.02e+04 | 5.85e+04 |
| x7 | 2.674e+04 | 2346.953 | 11.392 | 0.000 | 2.21e+04 | 3.13e+04 |
| x8 | 1.004e+05 | 2057.200 | 48.782 | 0.000 | 9.63e+04 | 1.04e+05 |
| x9 | 75.4452 | 1.989 | 37.939 | 0.000 | 71.547 | 79.343 |
| x10 | 38.3951 | 2.400 | 16.001 | 0.000 | 33.692 | 43.098 |
| x11 | -2605.2632 | 70.616 | -36.894 | 0.000 | -2743.675 | -2466.851 |
| x12 | 19.7064 | 3.641 | 5.412 | 0.000 | 12.569 | 26.844 |
| x13 | -598.1506 | 32.625 | -18.334 | 0.000 | -662.097 | -534.204 |
| x14 | 6.125e+05 | 1.06e+04 | 57.520 | 0.000 | 5.92e+05 | 6.33e+05 |
| x15 | -2.137e+05 | 1.26e+04 | -16.908 | 0.000 | -2.39e+05 | -1.89e+05 |

```

=====
Omnibus:                  18154.830    Durbin-Watson:              1.991
Prob(Omnibus):             0.000    Jarque-Bera (JB):          1781798.770
Skew:                      3.503    Prob(JB):                  0.00
Kurtosis:                  46.926    Cond. No.                  7.65e+16
=====

```

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified

```
[2] The smallest eigenvalue is 3.56e-20. This might indicate that there are
strong multicollinearity problems or that the design matrix is singular.
"""
```

Removed 'floors' and 'waterfront' 70% accuracy